

Rishikesh Reddy Regatte

rishikeshregatter@gmail.com | linkedin.com/in/rishikeshregatter/ | github.com/rishikesh-20

EDUCATION

University of North Carolina at Charlotte

Master of Science in Computer Science – GPA: 4.0

Hyderabad Institute of Technology and Management

Bachelor of Technology in Computer Science – GPA: 3.6

North Carolina, United States

January 2025 – December 2026

Telangana, India

August 2020 – August 2024

TECHNICAL SKILLS

Languages: Python, SQL (PostgreSQL), Java, C++, JavaScript, HTML/CSS

AI & Machine Learning: XGBoost, Scikit-learn, LangChain, Retrieval-Augmented Generation (RAG), Prompt Engineering, Pandas, NumPy, Matplotlib

Data Engineering: Databricks, Delta Lake, Apache Airflow, Airbyte, DuckDB, ETL Pipelines, Data Transformation, Data Warehousing, Data Modeling

Cloud & MLOps: AWS, MLflow, Unity Catalog, Git

Frameworks & Tools: React, Streamlit, Node.js, REST APIs, Power BI, Power Query, Excel

EXPERIENCE

Kastech Technologies

Telangana, India

Data Engineering Intern

August 2024 – November 2024

- Modeled and ingested attendance data into structured schemas to power Metabase dashboards, enabling dynamic HR reporting with real-time workforce visibility.
- Supported monthly financial reporting for 5 client companies by writing SQL queries to extract and prepare P&L and balance sheet data for senior review and downstream consumers.
- Optimized ETL pipelines by replacing full-refresh loads with incremental ingestion, cutting runtime from 10 to under 5 minutes for attendance and client financial data pipelines into PostgreSQL.
- Orchestrated client data ingestion workflows using Apache Airflow and Airbyte, and refactored Python transformation scripts to standardize data quality checks, saving 3+ hours of manual data preparation per week across reporting pipelines.

PROJECTS

LLM Prompt Benchmark | Python, Qwen3-14B, Gemini, RAGBench, BigQuery

- Built a local LLM evaluation harness benchmarking 5 prompting strategies across 150 RAGBench samples (750 Qwen3-14B generations), scored by a Gemini RAGAS-style judge for faithfulness, relevancy, and hallucination rate.
- RAG reduced hallucination from 68.7% to 10.5% vs. zero-shot; structured RAG reached 91.2% faithfulness while zero-shot scored 86% relevancy at only 31.3% faithfulness, revealing grounding and on-topic quality as independent axes.
- Engineered a resumable skip-logic runner with atomic JSON checkpoints, Streamlit dashboard, and optional GCP export to BigQuery and Cloud Storage for downstream reporting.

SEC Filing Intelligence Lakehouse | Databricks, Delta Lake, SQL, XGBoost, MLflow

- Architected a production-style Lakehouse and MLOps platform on Databricks to ingest, transform, and model SEC financial data for S&P 500 companies using Delta Lake, Unity Catalog, and Databricks Asset Bundles.
- Designed leakage-safe feature engineering pipelines and optimized a calibrated XGBoost model achieving 0.663 ROC-AUC and 62% accuracy through chronological validation and MLflow-driven experiment tracking.
- Operationalized end-to-end model deployment with automated batch scoring, champion model versioning, AI/BI dashboards, and peer clustering (Silhouette Score: 0.75) to support scalable financial intelligence.

Conversational Deductive Query Assistant | Python, SQL, Streamlit

- Developed a natural language to SQL application enabling dynamic querying of relational databases using AI-generated queries.
- Improved query generation accuracy by 35% through schema validation and structured prompt refinement techniques.
- Built interactive Streamlit dashboards providing tabular and visual outputs to accelerate exploratory data analysis workflows.

Medbot | LLM, Python, LangChain, Pinecone

- Designed and developed an AI-powered medical chatbot delivering context-aware health information using LangChain with Grok API and Retrieval-Augmented Generation (RAG) techniques.
- Integrated Pinecone vector database for semantic search, enhancing retrieval accuracy and reducing irrelevant responses by over 40%.
- Developed an interactive Streamlit-based user interface with secure API integration, enabling real-time querying and efficient information delivery.

CERTIFICATIONS

Anthropic: Claude Code in Action — Anthropic, March 2026

AWS: Data Engineer Certification — AWS, December 2025

IBM: Python for Data Science — IBM, August 2024

IBM: Machine Learning with Python — IBM, August 2024